

SRIKANTH

Rockville, MD | [+1 \(252\) 284-5365](tel:+12522845365) | [Email](#) | [Website](#) | [LinkedIn](#) | [GitHub](#)

SUMMARY

Software and data engineer with about 2 years of experience at redBus (India's largest online bus ticketing platform) and an MS from Wilmington University. Built production features in Python, ran performance testing on APIs, and worked on AI snippet integration for the redBus app and website. Master's capstone: a streaming CDC pipeline using Debezium, Kafka, Flink, and Delta Lake. Looking for software engineering or data engineering roles.

SKILLS

Languages: Python, SQL, C

Data & Streaming: Apache Kafka, Apache Flink, Debezium, Delta Lake, Apache Airflow, dbt, Pandas, Great Expectations

Backend & Tools: FastAPI, SQLAlchemy, Pydantic, PostgreSQL, REST APIs, Docker, Git, Grafana

Practices: Data pipelines (batch & streaming), ETL/ELT, CDC, data modeling, data quality, performance testing

EXPERIENCE

Software Engineer — redBus, Bengaluru, India

Feb 2022 – Jul 2023

- Built new features for the redBus app and website in Python; fixed production bugs across the platform.
- Worked on AI snippet integration for the app and website, adding LLM-generated content into the live product.
- Wrote Python scripts and workflows to move, clean, and prepare data for analytics.
- Ran performance and load tests on APIs and services to find slow endpoints and bottlenecks.

Software Engineering Intern — redBus, Bengaluru, India

Aug 2021 – Jan 2022

- Wrote technical documentation for platform components, mapping system architecture and data flows.
- Ran load tests on APIs to flag performance issues; converted to full-time after the internship.

PROJECTS

Streaming CDC Platform with Data Quality Enforcement

Master's Capstone (Grade: A)

Python, PostgreSQL, Debezium, Kafka, Flink, Delta Lake, Great Expectations, Grafana, Docker

- Built a near-real-time pipeline syncing a Postgres database to an analytics layer without polling the source.
- Captured row-level changes from the write-ahead log with Debezium and Kafka; processed the stream in Flink, handling deduplication, late-arriving events, and schema evolution.
- Merged data into Delta Lake tables via upserts; Great Expectations quality gates quarantined bad records instead of breaking the pipeline.
- Set up a Grafana dashboard for consumer lag, throughput, and data freshness; wrote a Python transaction simulator (orders, users, payments) as the source app.

DriveLingo AI — Traffic Sign Intelligence Pipeline

Personal Project

Python, FastAPI, PostgreSQL, SQLAlchemy, Pydantic, OCR

- Building a system that turns traffic signs a camera sees — in any language — into what the sign says, in the driver's language, and why it matters.
- Built so far: ingestion storing raw events separately from processed records, Pydantic validation, database-level dedup, and an OCR worker extracting sign text from images.
- Roadmap: language detection and translation, AI-generated safety explanations, driver alerts, and analytics on OCR confidence and sign frequency by region.

LocalLoop — Local Vendor Intelligence Platform

Personal Project (in progress)

Python, FastAPI, PostgreSQL, dbt, Airflow, Pandas, Docker

- Building an event-driven platform for re-finding trusted local vendors via NFC tap or QR scan after a job.
- Anonymous interaction events feed vendor-retention and locality-level analytics with dbt and Airflow; customer identity stays private and vendors see only aggregated metrics.

EDUCATION

Wilmington University — MS, Information Systems and Technology

Sep 2023 – May 2025

Concentration: Information Assurance | New Castle, DE

Key courses: Data Management, Cloud Management, Networking, Web and Data Security.

JNTU Hyderabad — B.Tech, Electronics and Communication Engineering

2018 – 2022

Final year project (team): RFID-based student campus exit tracking using Arduino and C.